



VR Direct: How NVIDIA Technology Is Improving the VR Experience

Nathan Reed — Developer Technology Engineer, NVIDIA

Dario Sancho — Lead Programmer, Crytek

Who We Are

- Nathan Reed
 - NVIDIA DevTech — 2 yrs
 - Previously: game graphics programmer at Sucker Punch
- Dario Sancho
 - Crytek — 2 ½ yrs
 - Previously: academy, system and platform programming



Hard Problems of VR

Headset design


Input

Rendering performance


Experience design



Latency



Stereo Rendering



Two eyes, same scene



What Is VR Direct?

- Various NV hardware & software technologies
- Targeted at VR rendering performance
 - Reduce latency
 - Accelerate stereo rendering



VR Direct Components

In This Talk

Asynchronous Timewarp

VR SLI



Latency

Frame Queuing

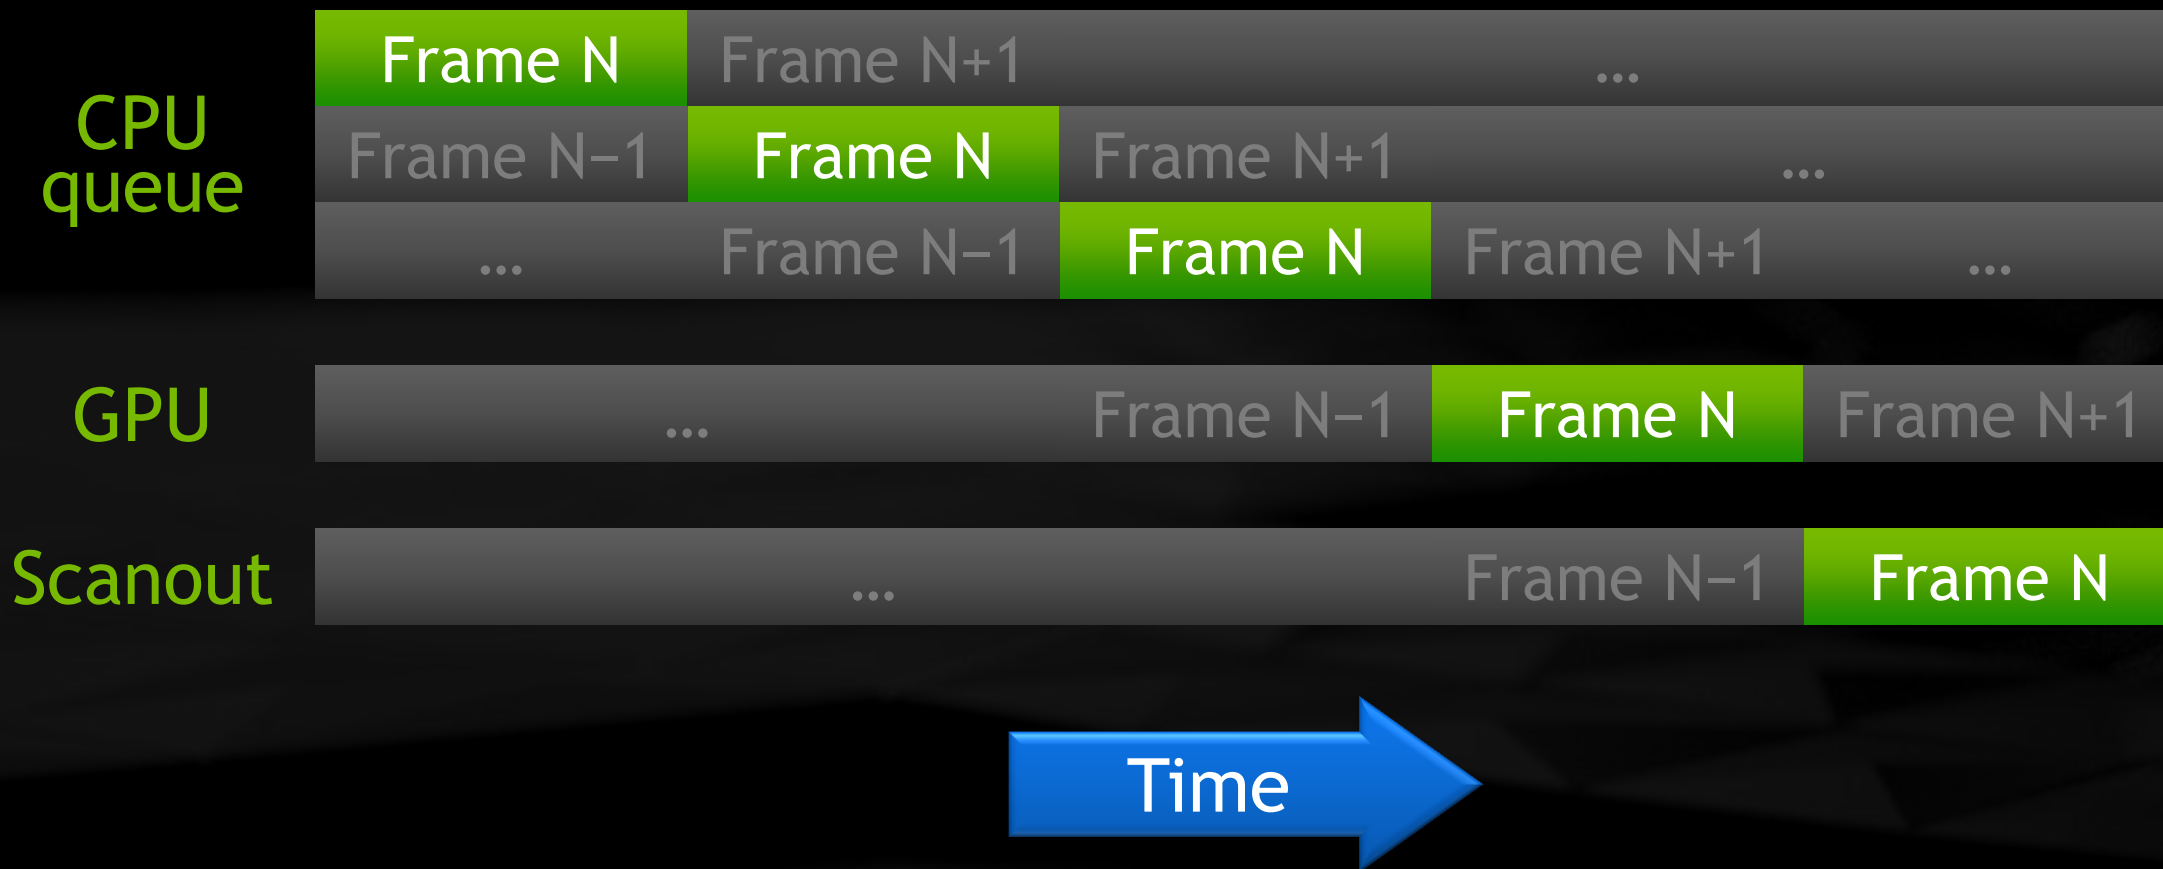
Timewarp

Late-Latching Constants

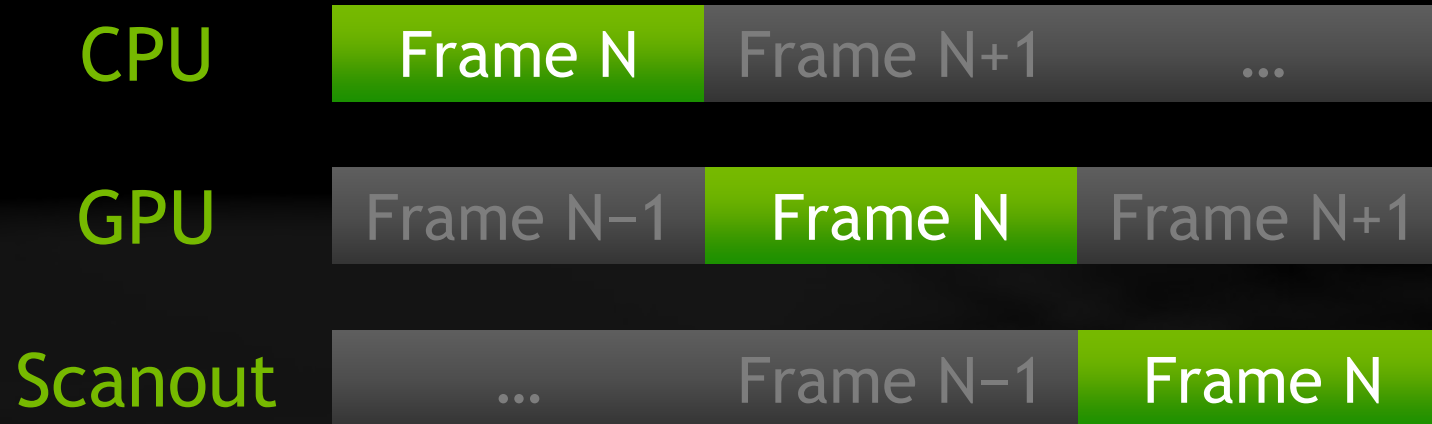
Asynchronous Timewarp





Frame Queuing



Frame Queuing



Timewarp

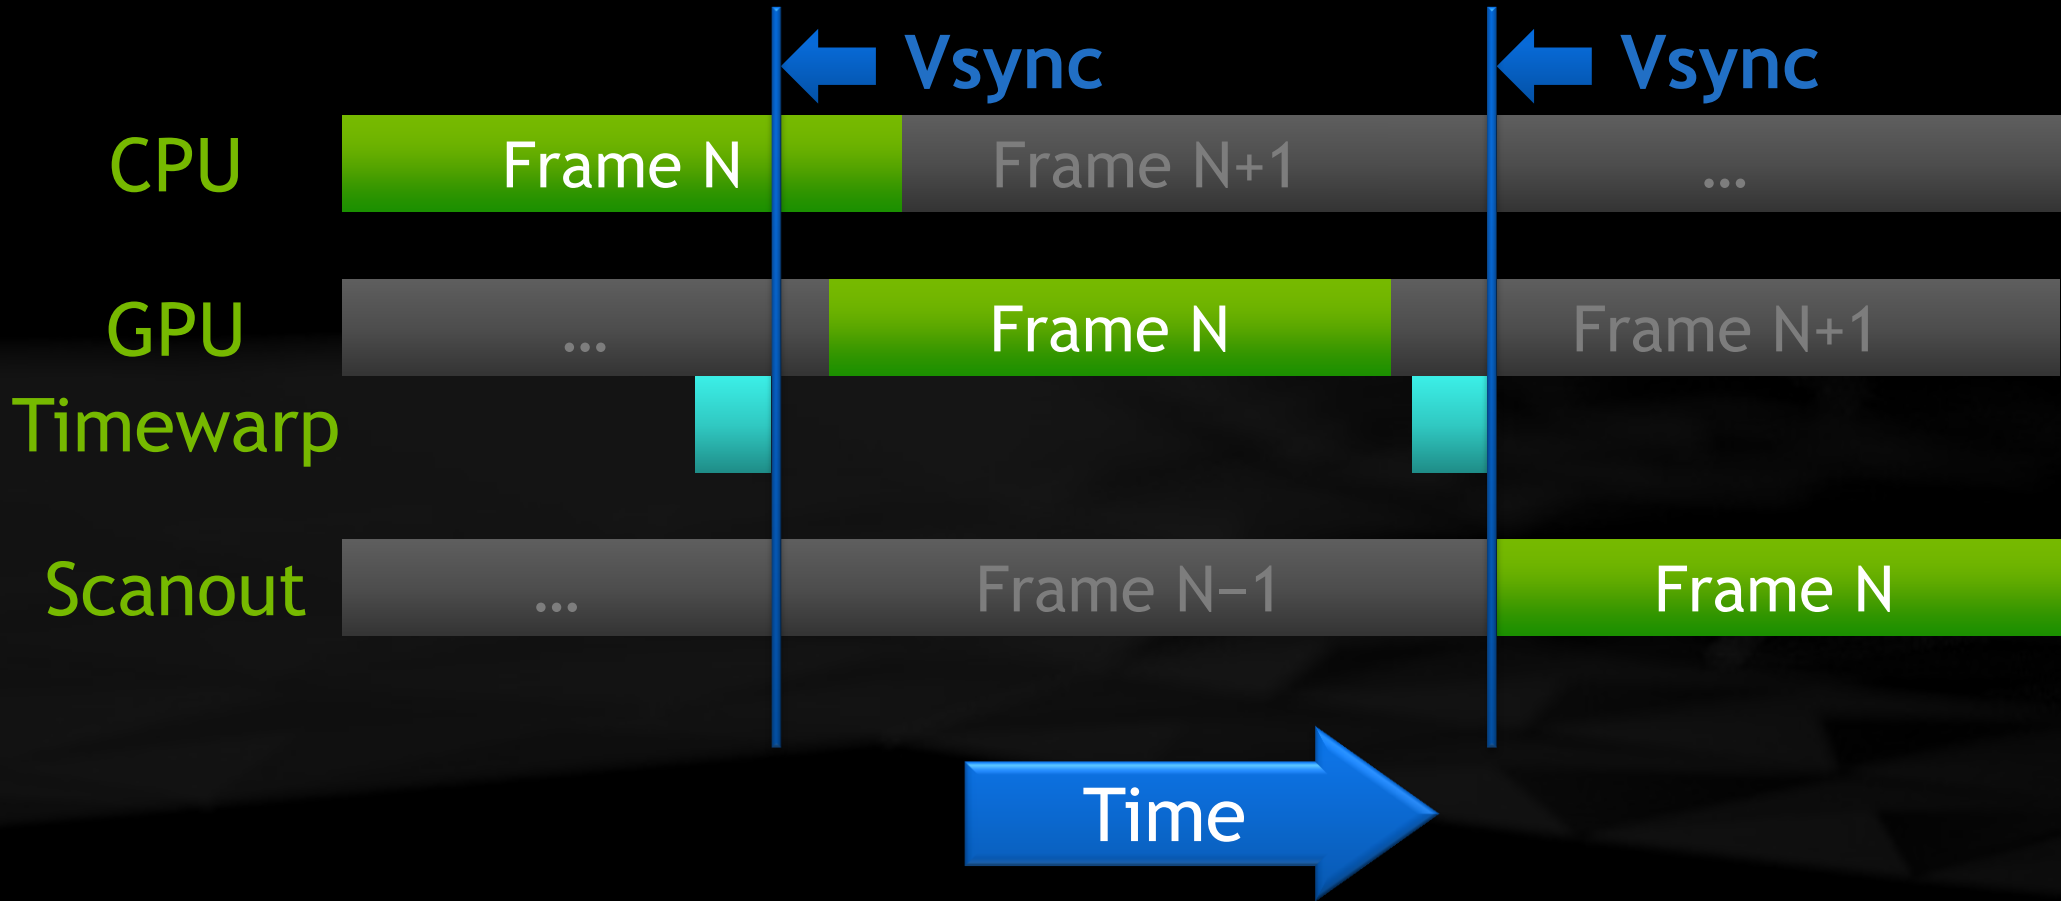


Timewarp Pros & Cons


- Very effective at reducing latency...of rotation!
 - Fortunately, that's the most important
- Doesn't help translation!
- Doesn't help other input latency
- Doesn't help if vsync is missed



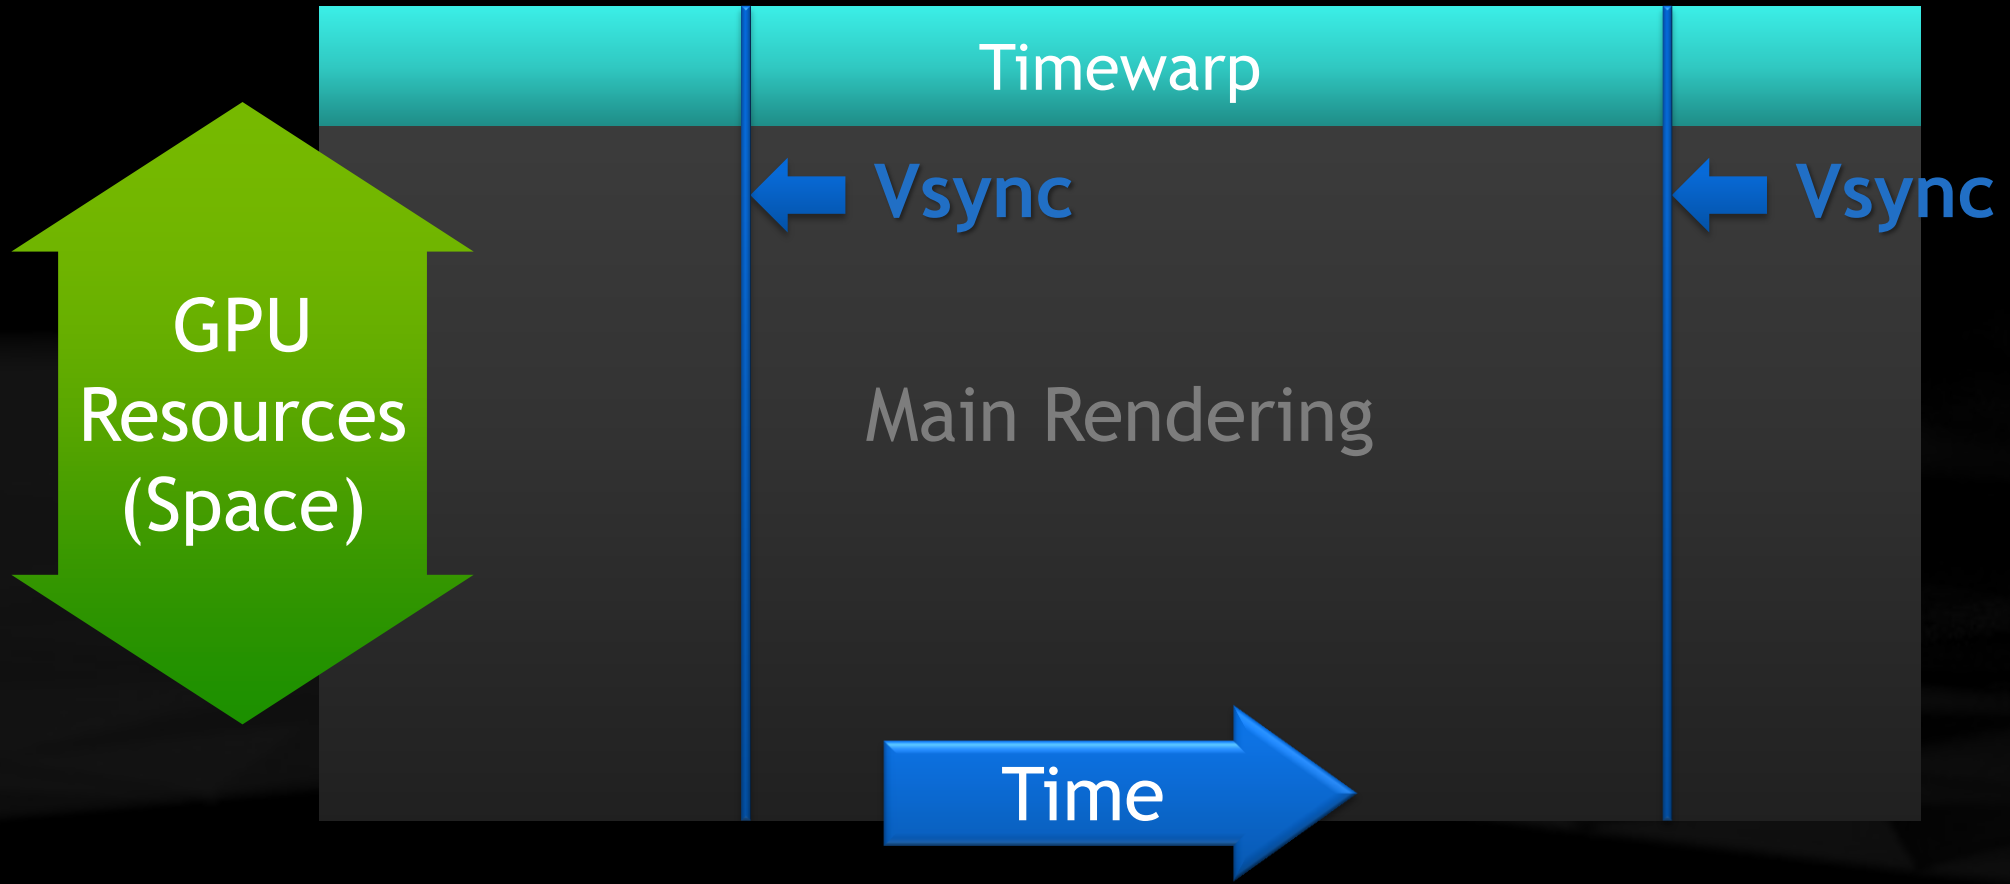
Asynchronous Timewarp



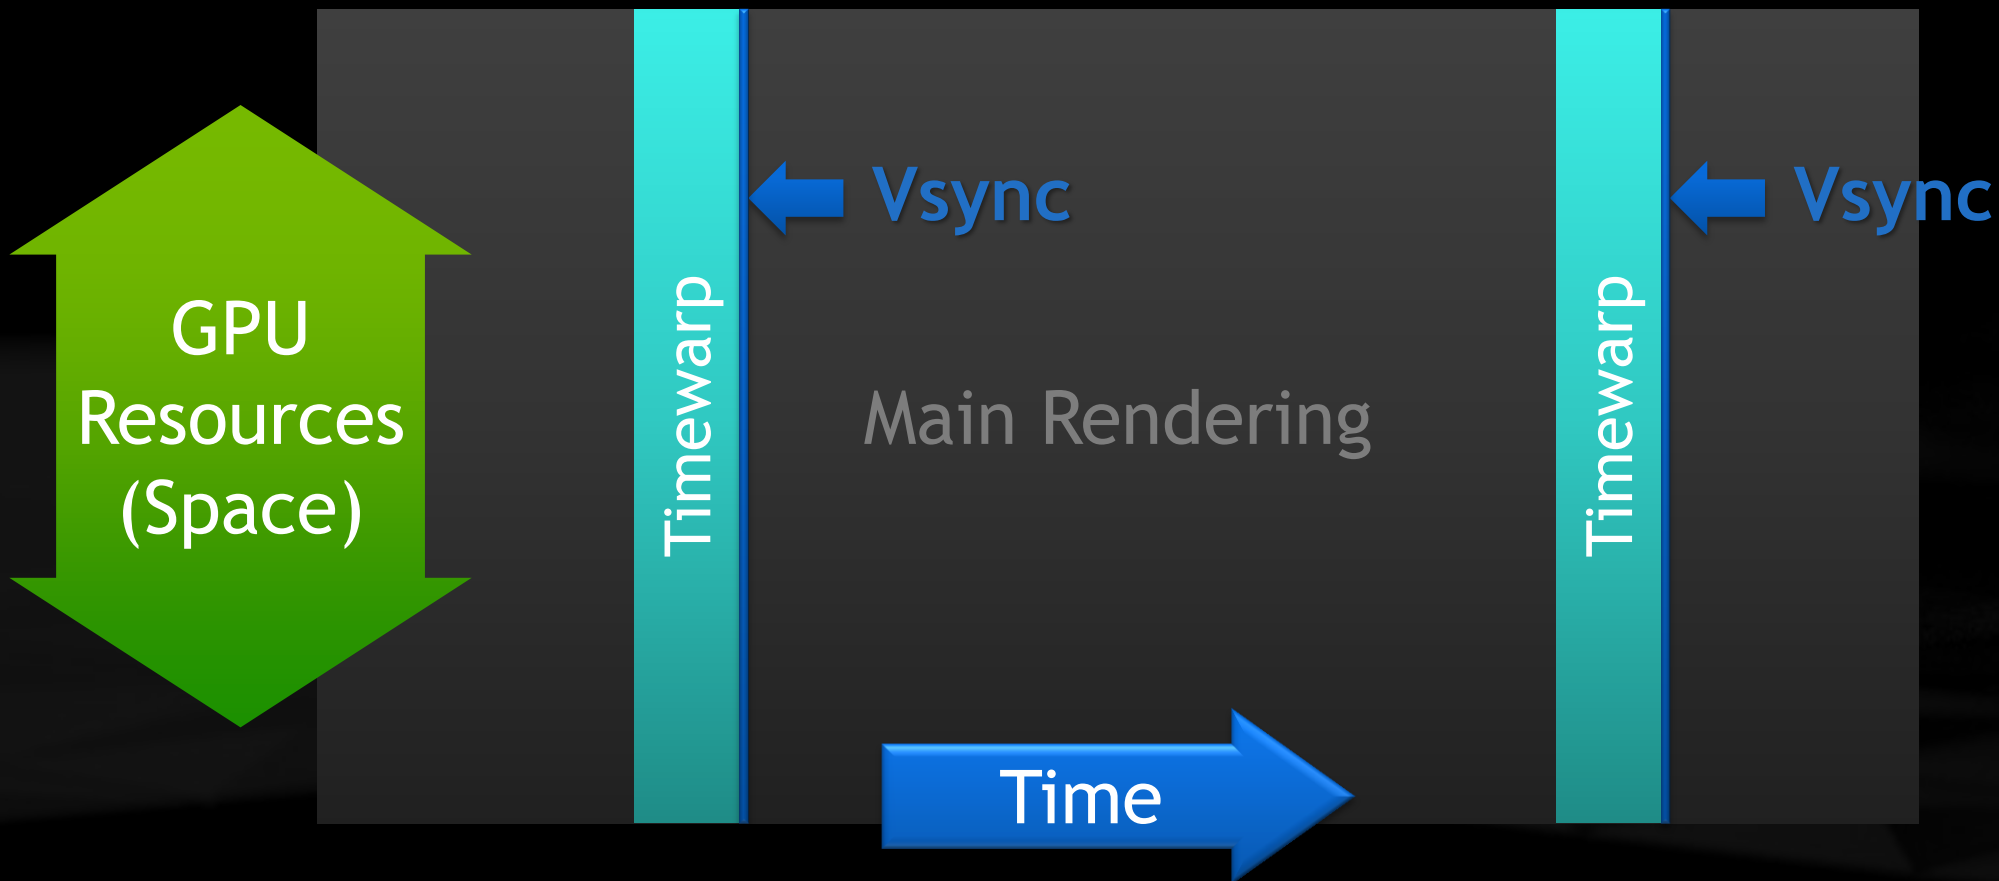
Space Vs Time



Space-Multiplexing



Time-Multiplexing

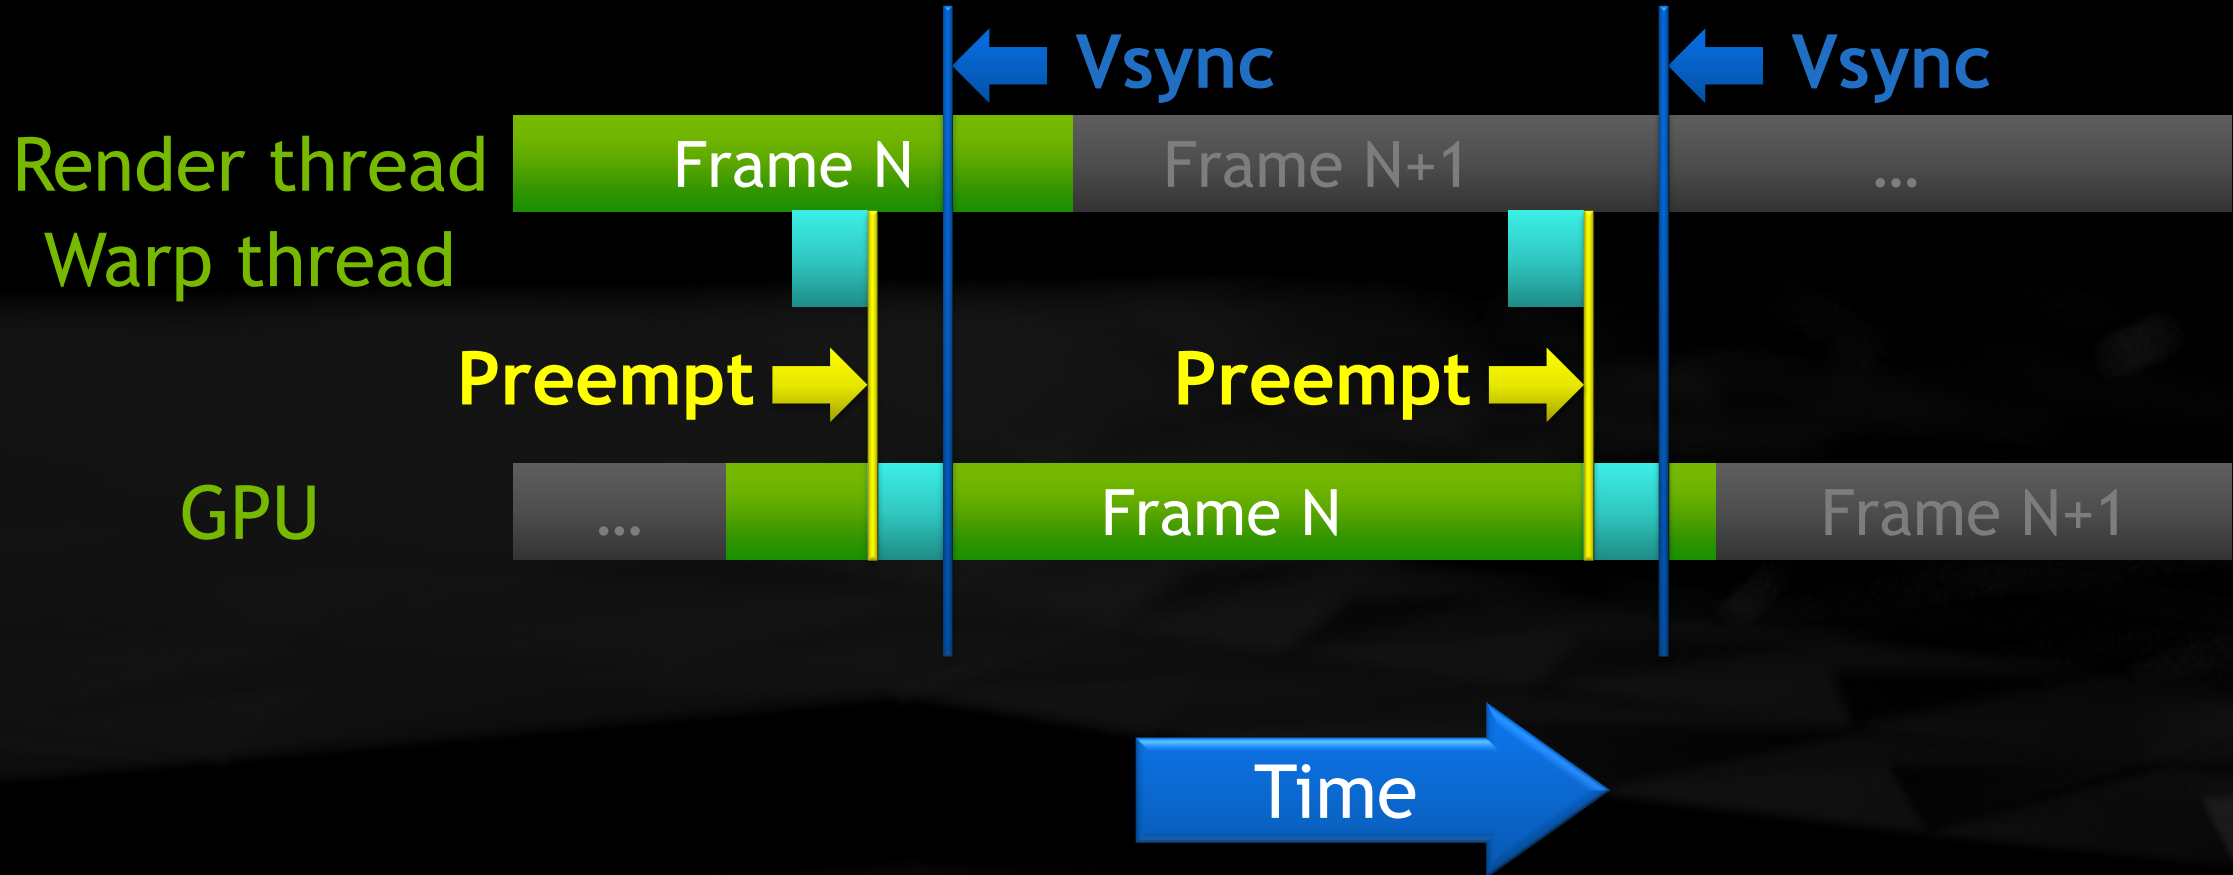


High-Priority Context

- NV driver supports high-priority graphics context
 - Time-multiplexed – takes over entire GPU
- Main rendering → normal context
- Timewarp rendering → high-pri context



Async Timewarp With High-Pri Context



Preemption

- Fermi, Kepler, Maxwell: draw-level preemption
- Can only switch at draw call boundaries!
 - Long draw will delay context switch
- Future GPU: finer-grained preemption



Direct3D High-Priority Context

- `NvAPI_D3D1x_HintCreateLowLatencyDevice()`
- Applies to next D3D device created
- Fermi, Kepler, Maxwell / Windows Vista+
- NDA developer driver available now



OpenGL High-Priority Context

- `EGL_IMG_context_priority`
- Adds priority attribute to `eglCreateContext`
- Available on Tegra K1, X1
 - Including SHIELD console
- Only for EGL (Android) at present
 - WGL (Windows), GLX (Linux) to come



Developer Guidance

- Still try to render at headset native framerate!
- Async timewarp is a safety net
 - Hide occasional hitches / perf drops
 - Not for upsampling framerate



Developer Guidance

- Avoid long draw calls
 - Current GPUs only preempt at draw call boundaries
 - Async timewarp can get stuck behind long draws
- Split up draws that take >1 ms or so
 - E.g. heavy postprocessing
 - Split into screen-space tiles



Latency TL;DR

- Reduce queued frames to 1
- Timewarp: adjusts rendered image for late head rotation
- Async timewarp: safety net for missed vsync
- NVIDIA enables async timewarp via high-pri context



Stereo Rendering

Multiview Rendering

VR SLI



Frame Pipeline

Which stages must be done twice for stereo?

CPU

Find visible objects



Submit render commands



Driver internal work

GPU

Transform geometry



Rasterization




Shading




Flexibility vs Optimizability

More flexible — all stages separate



Flexibility vs Optimizability

More optimizable — some stages shared



Stereo Views

- Almost the same visible objects
- Almost the same render commands
- Almost the same driver internal work
- Almost the same geometry rendered



Other Multi-View Scenarios

- Cubemaps: 6 faces
- Shadow maps
 - Several lights in one scene
 - Slices of a cascaded shadow map
- Light probes for GI
 - Many probe positions in one scene

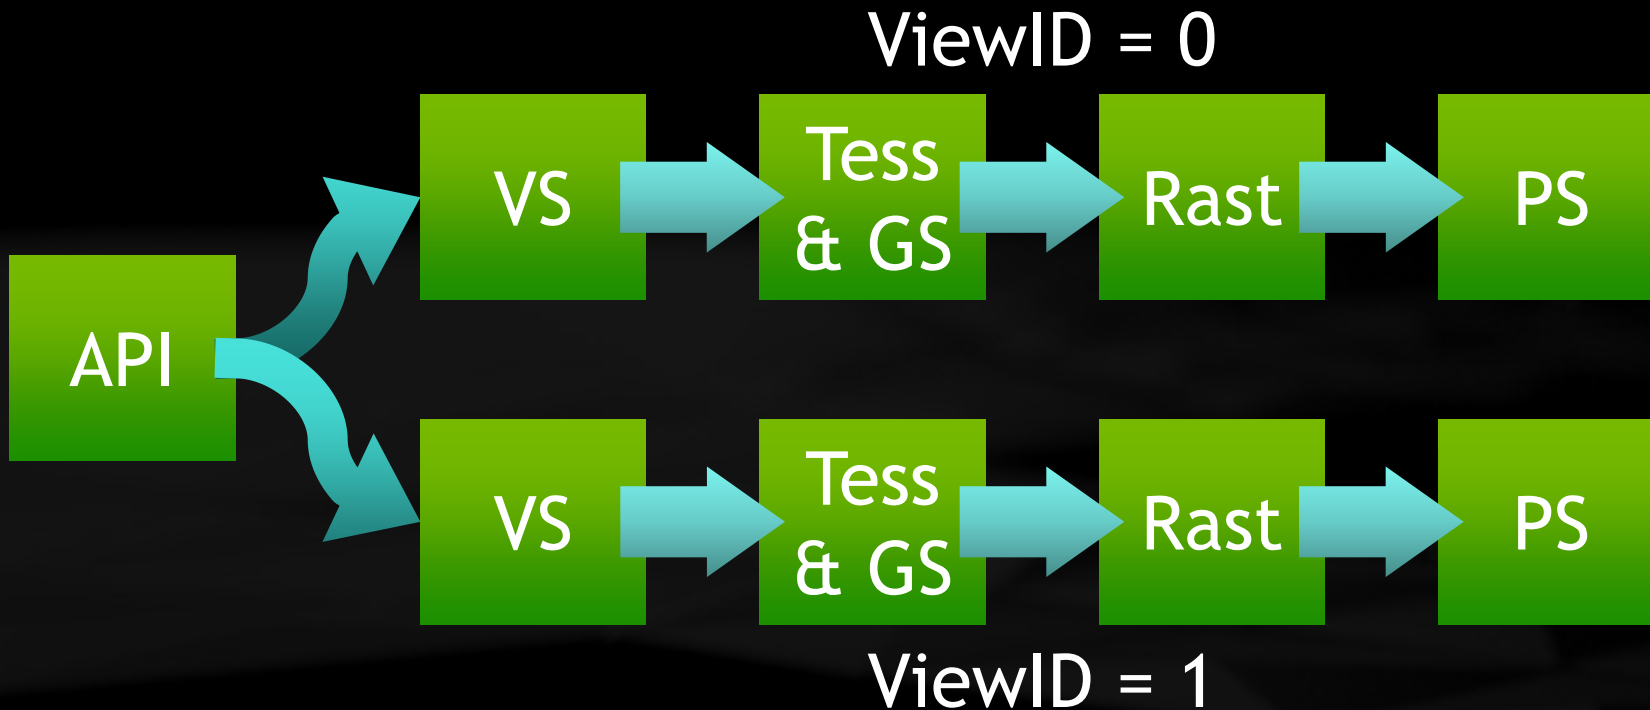


Multiview Rendering


- Submit scene render commands *once*
- All draws, states, etc. broadcast to all views
- API support for limited per-view state
- Saves CPU rendering cost
- Maybe GPU too — depending on impl!



Shader Multiview




Hardware Multiview




VR SLI

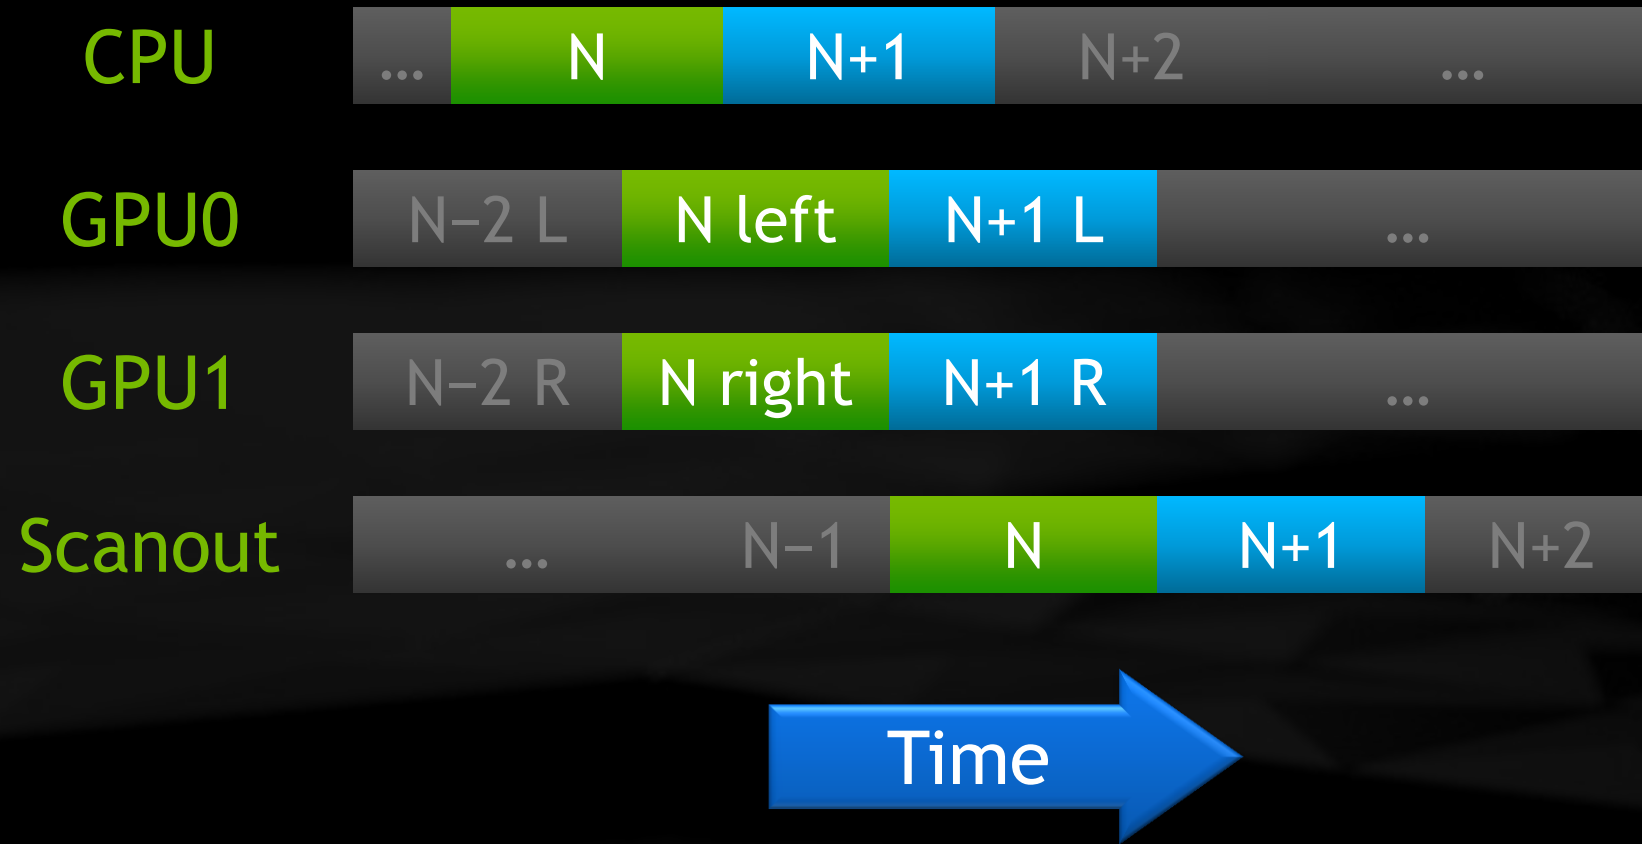
Shared
command
stream



Interlude: AFR SLI



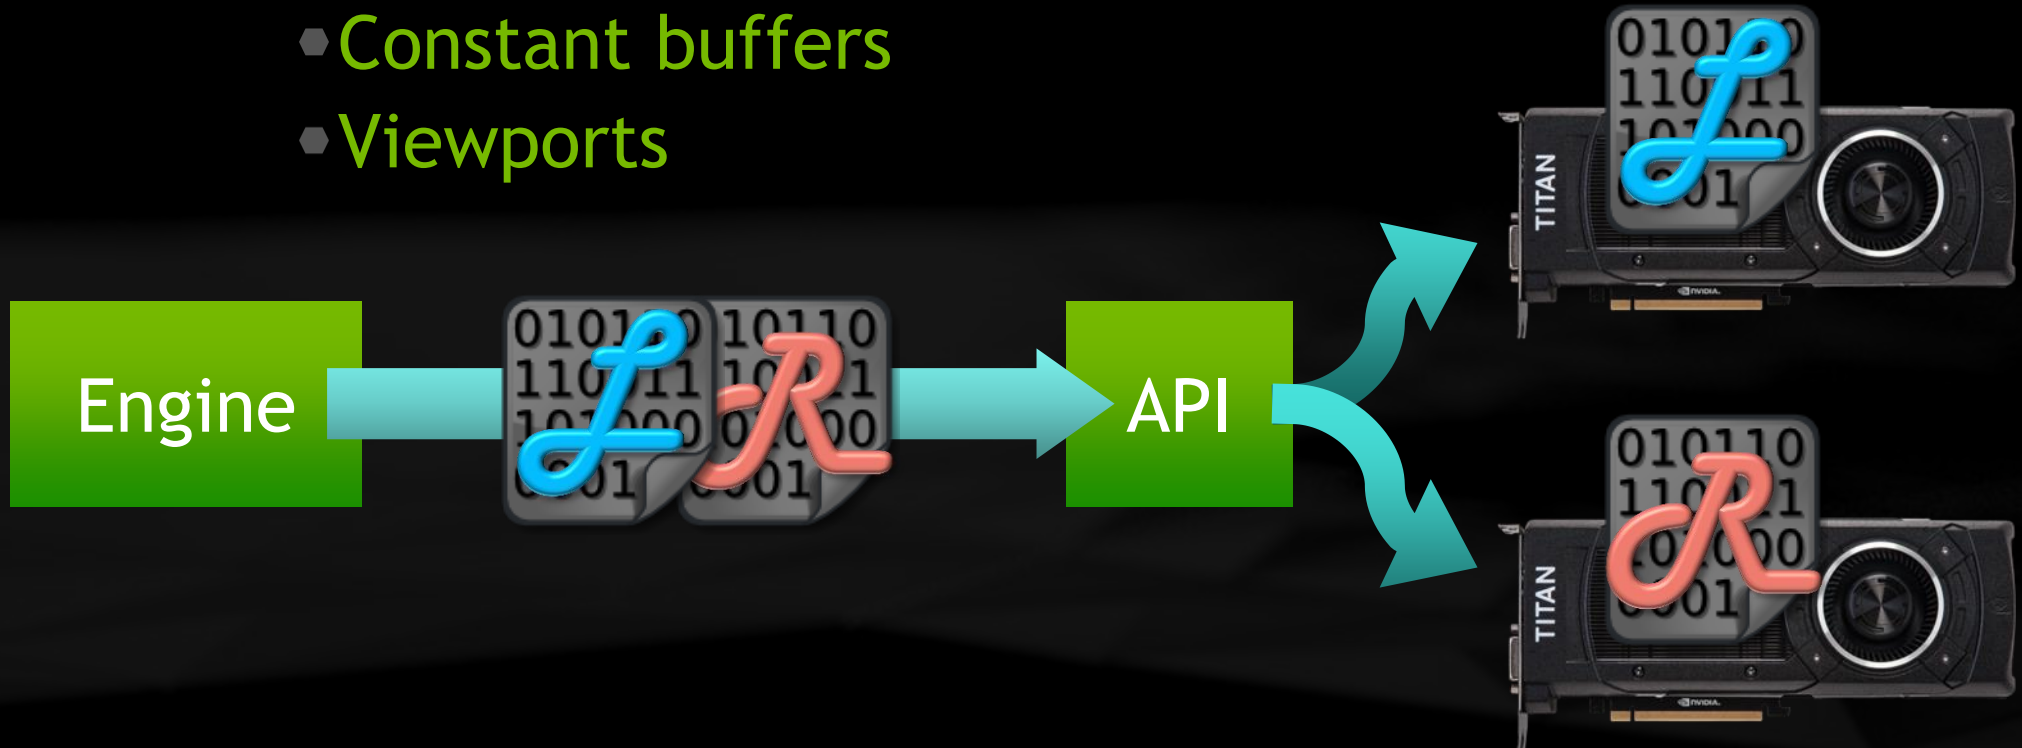
VR SLI



VR SLI

Per-GPU state:

- Constant buffers
- Viewports



VR SLI



VR SLI Scaling

- View-independent work (e.g. shadow maps) is duplicated
- Scaling depends on proportion of view-dependent work



API Availability

- Currently D3D11 only
- Fermi, Kepler, Maxwell / Windows 7+
- Developer driver available now
- OpenGL and other APIs: to come



Developer Guidance

- Teach your engine the concept of a “multiview set”
 - Related views that will be rendered together

- Currently:

```
for (each view)
    find_objects();
for (each object)
    update_constants();
render();
```



Developer Guidance

- **Multiview:**

```
find_objects();  
for (each object)  
    for (each view)  
        update_constants();  
render();
```



Developer Guidance

- Keep track of which render targets store stereo data
 - May need to be marked or set up specially
 - Or allocated as a texture array, etc.
- Keep track of sync points
 - Where you need all views finished before continuing
 - May need to blit between GPUs



Stereo Rendering TL;DR

- Multiview: submit scene once, save CPU overhead
 - Requires some engine integration
- Range of possible implementations
 - Trade off flexibility vs optimizability
- VR SLI: a GPU per eye



VR Direct Recap

- Variety of VR-related APIs coming in near future
- Reduce latency
 - Reduced frame queuing
 - Enable async timewarp & other improvements
- Accelerate stereo rendering
 - Multiview APIs
 - VR SLI



VR Direct API Availability

- Fermi, Kepler, Maxwell
- D3D11: context priorities and VR SLI
 - NDA developer driver available now
- Android: EGL_IMG_context_priority
- Other APIs/platforms: to come



What Next?

- All this stuff is hot out of the oven!
- Will need more iterations before it settles
 - See what works, revise APIs as needed
 - Consolidate & standardize across industry



How VR Is Shaping CryEngine



Our VR Challenges

• As Developer

- Focus on results: Best possible VR demo for GDC 2015 (presence, interaction, performance...)
- Focus on the platform to be shown
- Short development time

• As Technology Provider

- Solid implementation
- Multiplatform and support for multiple head set vendors
- Focus on performance and seamless integration for users



CRYENGINE®



Exploring the key aspects of VR

- Presence
 - Convincing rich environments, 3D audio, etc.
- Interactivity
 - Allow the player to manipulate the world instead of just watching
- Input devices
 - Experimenting with traditional and next-gen input devices
- Movement
 - Believable, stable, non-sickening

Our Rendering Challenges

- **High & stable frame rate**
 - Oculus requires 90+ FPS (drops → physical discomfort)
- **Resolution:** Full HD and beyond
- **Quality:** Bringing our signature visuals to VR
- **Dual rendering vs** Reprojection
- **Minimum latency**

VR Demo Rendering Tale

REPROJECTION



GPU

Single GPU solution – 90+ FPS at full HD



3D TV

Excellent performance, just a post-process
Worked well on 3D TVs



VR
(R&D)

Introduces artefacts
Reduced depth perception (to minimize artefacts)
→ “Presence” is not fully achieved



VR Demo Rendering Tale

REPROJECTION



DUAL RENDERING



MGPU



AFR

VR SLI
(engine
pipeline)



GPU Masking

(good proof of concept)

Single GPU (brute force)



CPU, GPU: 2x Rendering Effort!!



3D TV



VR
(R&D)

VR Demo Rendering Tale


REPROJECTION



DUAL RENDERING



MGPU



No driver support



Single GPU (brute force)



CPU, GPU: 2x Rendering Effort!!


Solution: Use Next-gen GPU



Titan X (**GM200 GPU**)

Engine Optimization
Art optimization

90+ FPS (around 100 FPS avg.)

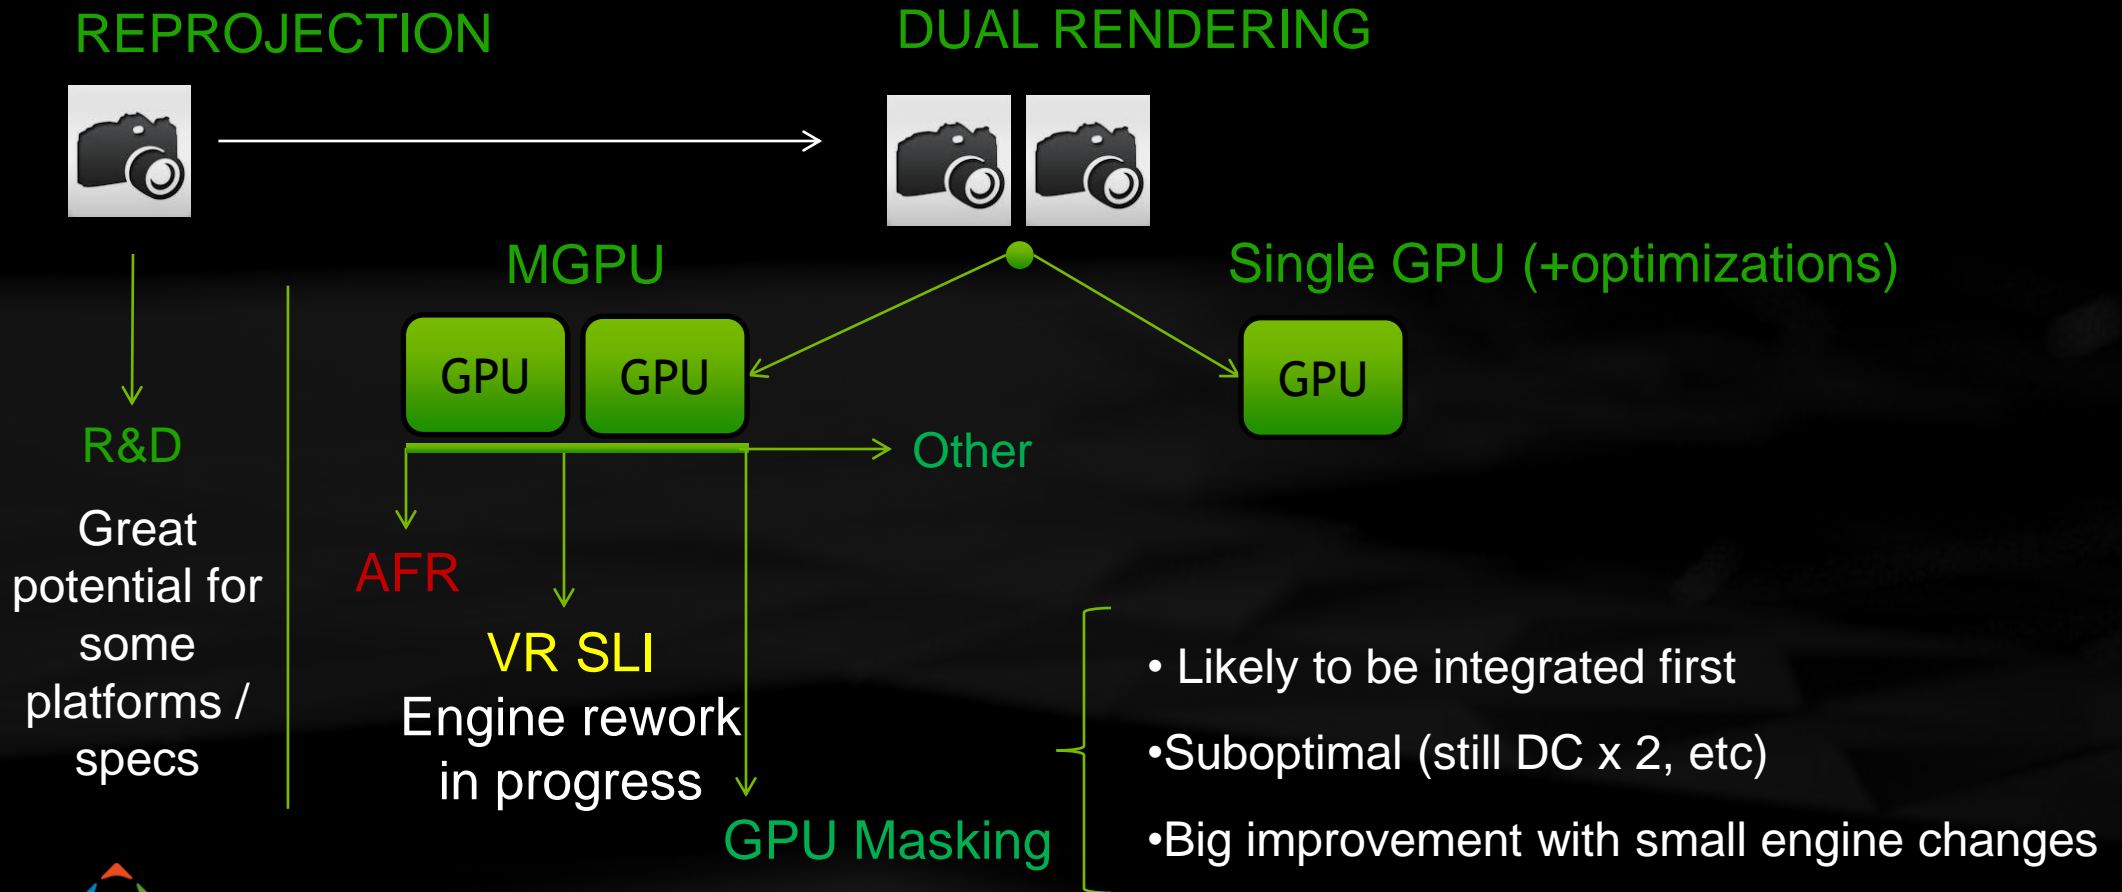


3D TV



VR
(R&D)

CryEngine VR Integration



CryEngine VR Integration

REPROJECTION



DUAL RENDERING



MGPU

Single GPU



R&D

Great
potential for
some
platforms /
specs

Not just a tech demo

Community wants to work with these technologies ASAP
Licensees and potential partners' requests are the proof

Questions & Comments?

Email us:

nreed@nvidia.com

dariop@crytek.com

Slides will be posted:

<http://gputechconf.com>

